
Decodable Is Not Learned: Untrained-Weights Controls and Basis-Starved Splits for Linear Probing

Michael Jordan
Independent researcher

michael.jordan@metrc.com

Abstract

A linear probe that clears significance is routinely read as evidence that a model has learned, represents, or is developing a capability. The inference is unsound: a network’s activations are a high-dimensional expansion of the visible tokens, and a linear readout on such an expansion decodes many functions of the input that training never put there. This paper prescribes three controls that make the inference testable, an untrained-weights twin of the model run through the identical probe pipeline as an acceptance test, basis-starved splits whose validation items share no surface-component values with probe training, and tolerances calibrated from the mechanism that generates each control’s events, and reports two preregistered probing campaigns that the controls terminated before any outcome measurement. In the first, the untrained twin fired on 120 of 120 fits: under standard splits, probe significance measured reservoir decodability, not learning. In the second, starving closed the lookup class and the twin exposed a class beneath it, 13 of 25 capabilities decodable from untrained weights at margins of .06 to .82 via surface statistics that no value holdout can remove; the twelve survivors read exactly zero untrained margin in every cell. A successor battery built with the screen at inclusion time froze with the untrained gate clean, the program’s first, and then committed the outcome-side form of the same error: an untrained floor sits at zero on a generation task, so it credited format acquisition as capability, and re-scoring against within-answer-space chance moved the headline rank correlation from .368 to .200. The two capabilities carrying the second and third highest starved margins scored exactly zero argmax accuracy at every scale: decodable is not learned, and on this evidence not generable either. A taxonomy of the leak mechanisms and a checklist written for verbatim adoption distill the record.

1 The problem and the prescription

A linear probe reads a label off a network’s activations, clears a significance test, and the paper concludes the model represents the property. The step from “a readout can be fit” to “the model learned something” is taken silently in most probing work, and it isn’t sound. A transformer’s activations at any layer are, among other things, a high-dimensional nonlinear expansion of the visible tokens, and an expansion like that supports linear readouts of many functions of the input that no training ever put there. Random-feature regression and reservoir computing are whole method families built on exactly this capacity of untrained networks; Section 2 gives the arithmetic. The fraction of probe signal that comes free with the architecture is not small and not exotic: in the experiments reported here, an untrained network cleared the same significance tests as the trained one on every capability in a twelve-task battery.

The prescription is short enough to state completely on page one.

- **P1. The untrained twin as an acceptance test.** Run the identical probe pipeline (same architecture, tokenizer, prompts, extraction points, splits, and statistics) on a twin of the model

Portions of this paper’s text, analysis code, and figures were prepared with the assistance of a large language model (Anthropic’s Claude), used as a drafting and engineering tool under the author’s direction; ideas, claims, experimental designs, and adjudication rulings are human-sourced. The Disclosure of AI assistance section describes the full division of labor.

with weights at seeded random initialization. Capability-precursor claims stand only on the trained-minus-untrained gap. The screen belongs at inclusion time, before a capability enters the battery; the campaigns reported here ran it post hoc and spent both batteries learning where it belongs. Section 7 reports the battery then built with the screen at inclusion.

- **P2. Basis-starved splits.** Wherever the label could be looked up from surface components of the prompt, name those components against the tokenizer’s actual token inventory, then validate only on items whose every component value was held out of probe training, discarding items that mix held-out and kept values.
- **P3. Mechanism-calibrated tolerances.** Every control gets a tolerance derived from the mechanism that generates its events: no zero-tolerance rules on tests with designed nonzero rates, and per-fire signature bars from the order statistics of the permutation null (the expected maximum of N permuted fits), not from SD intuitions.

The evidence is two preregistered probing campaigns, both ended by these controls before either touched its outcome variable. In the first, the untrained control fired on 120 of 120 fits, each at the permutation floor, across every capability: standard train-validation splits let a linear readout assemble lookup tables in the random expansion, and diagnostics showed the trained probes’ ceiling margins were lookup too. In the second, splits starved by design closed that class (the prior campaign’s proven offender fell from margin 1.0 to about 0.1 untrained), and the untrained control then exposed a second class underneath: 86 of 250 untrained fits fired structurally, on thirteen of twenty-five capabilities, at margins from .06 to .82 that reproduce across two independently initialized models — labels partially computable from surface statistics that no holdout of basis values can remove. Both experiments returned `INSUFFICIENT_DATA` under their frozen rules. The deaths are the point: every failure was a detection, and every verdict was projected in a timestamped ledger before the frozen adjudication ran.

The paper contributes, in order of expected reuse, the following. An operational checklist for probing hygiene, written for verbatim adoption (Table T2, Section 10). A taxonomy of surface-computability leaks with per-capability anatomy, separating rescuable tasks from surface-all-the-way-down ones (Section 5). The starved-probe instrument with its calibrated gates and determinism infrastructure (Section 3), and the two campaign records as worked examples of the controls under preregistration discipline (Section 4). The successor battery’s screening arc, in which the controls ran at inclusion time, confirmed three taxonomy mechanisms by construction and produced the program’s first clean untrained gate (Section 7), and the outcome-side error that battery then committed (Section 8). And eleven calibration rules for frozen criteria, learned from defects in my own frozen code and promoted to standing practice, one of them measured by re-fitting what the campaign’s own records had discarded: a probe’s best site is a selected value, inflated 1.28x on untrained networks by the selection alone (Section 6).

2 Background

Linear probing entered the field as a diagnostic: fit a linear classifier on frozen intermediate activations and read its accuracy as evidence about what the network represents (Alain & Bengio, 2016). The trouble with that reading was noticed early, and the critique literature that followed is, in most of its branches, a series of constraints on the probe. Hewitt & Liang (2019) showed that an expressive probe can score well by learning the task itself rather than reading it out, and proposed control tasks: the same probe fit against randomized word-to-label mappings, with selectivity, the gap between task and control accuracy, as the honest report. Voita & Titov (2020) replaced accuracy with description length, separating a probe that finds structure cheaply from one that grinds it out of raw material. Elazar, Ravfogel, Jacovi, and Goldberg (2020) moved the question from decodability to use: erase the property from the representation and watch whether the model’s behavior degrades. Belinkov’s (2021) survey organizes the field around this family of worries. What the branches share is their target. Control tasks bound what the probe could learn from anything; description length prices the probe’s effort; amnesic probing asks whether the model rather than the probe depends on the property. All of them interrogate the probe’s expressivity. None of them interrogates the substrate’s: what a linear readout can extract from this architecture before training has shaped a single weight.

Untrained networks as comparison points do appear in the literature. Zhang & Bowman (2018) reported that randomly initialized, frozen encoders retain much of the probing signal trained ones show, an effect that vanished only when the probes’ training data was cut. Wieting & Kiela (2019) found random-projection sentence encoders close enough to trained ones on standard classification benchmarks to recommend them as mandatory baselines. But the untrained network has lived as a row in a table: a number the headline result should exceed, compared informally, with nothing decided when the comparison comes out badly. The prescription in this paper is that row promoted to an acceptance test. The untrained twin runs the identical pipeline (same prompts, splits, candidate sweep, permutation null, and correction), its tolerated fire count is derived from the pipeline’s own false-positive arithmetic, and battery membership is decided by the outcome. The distinction from control tasks is worth stating exactly: a control task randomizes the labels and keeps the trained representation, measuring what the probe could do with anything; the untrained control keeps the true labels and randomizes the representation, measuring what this architecture gives away for free. Section 4 shows the two catch different failures, and that the informal version of the comparison saturates precisely where it’s needed most.

There’s nothing mysterious about a random network that computes; an engineering field is built on it. Reservoir computing treats a fixed random recurrent network as a high-dimensional expansion of its input stream and trains only a linear readout, on the explicit design premise that random dynamics furnish a feature basis for free (Jaeger, 2001; Maass et al., 2002). Rahimi & Recht (2007) supplied the feedforward theory: linear fits on random nonlinear projections approximate kernel machines. Cover (1965) supplied the arithmetic that says when to worry: a linear separator with d degrees of freedom has a separating capacity of about $2d$ points in general position. The probes in Sections 4 and 5 fit on the order of 1,500 training rows against hidden widths of 1,024 at 410M and 2,048 at 1B, at or under capacity, where a clean training fit is free and means nothing. Everything rests on validation, and validation is where the reservoir property bites: a readout generalizes whenever the label is a simple enough function of surface features that the random expansion preserves them. A probing study at these scales that skips the comparison is asserting that its accuracy came from learning while standing on a substrate the reservoir literature would use as shipped.

The batteries these controls were built for served a measurement program adjacent to the emergence debate. Wei et al. (2022) called an ability emergent when smaller models lack it and larger ones have it; Schaeffer, Miranda, and Koyejo (2023) argued that much of the apparent discontinuity is manufactured by discontinuous metrics. One way to press on that dispute is to look below threshold with something sharper than a benchmark: if linear-probe margins at 410M and 1B forecast which capabilities later appear up the Pythia ladder (Biderman et al., 2023), the capability was developing smoothly before the benchmark could see it. This paper argues neither side of that question; both campaigns closed before any outcome variable was measured. What the program contributed here is discipline: probe readings meant as forecasts had to be committed before the forecasted models were queried, so no outcome data existed to sanity-check the probes against, and the controls of Section 3 had to carry the whole burden of trust.

3 The instrument

The instrument is a linear probe with its evaluation moved onto ground the confound can’t reach. This section describes it as frozen for Experiment 2b; everything here was committed and tagged before the campaign produced data. The probe machinery is one experiment older than the campaigns of Section 4: Experiment 1 validated the instrument class on synthetic ground truth, small transformers trained from scratch on tasks where the presence or absence of the capability was known by construction, and the permutation-null probe module frozen there is the one Experiments 2 and 2b inherit. Its closeout also supplied the first entry in Section 6’s catalog of frozen-criterion defects, a magnitude criterion misspecified across incomparable chance floors, ledgered as such and left to stand.

Targets and items. Each capability contributes about 2,000 probe items: a prompt, the model’s hidden states over that prompt, and a probe label defined by the capability’s specification. The label is never the free-form answer but a small-alphabet function of it chosen at design time (the ones digit of the root, the first letter of the solution, the power of ten), which keeps probe classes balanced enough to test and forces every specification to say exactly what the probe is supposed to find. A specification must declare five things before

an item is generated: the probe target, the surface basis, the starving split with a feasibility count, an oracle, and a dumbest-baseline analysis stating what a lookup table or a random network should score. Capabilities that can't fill those fields honestly don't enter the battery. Table 1 lists the scored battery in full.

The probe. Activations are harvested at two fixed token positions per layer (position 0 is the question-end token), with layers thinned to every third plus the final one: nine layer positions at 410M and seven at 1B, so a capability's fit sweeps a family of 18 or 14 (layer, position) candidates. Each candidate is a scaler and a logistic regression ($C = 1.0$, 100 iterations) fit on probe training rows and scored as accuracy on the starved validation rows, with the scaler fit inside the split. Significance comes from a permutation null of 2,500 label permutations over all items under the fixed split, refitting the probe each time; under the null of no activation-label association this is exact regardless of the split's structure, because a lookup strategy is starved in the observed fit and in every permuted fit alike. The best candidate's p-value is Bonferroni corrected for the family, a capability reads "present" when the corrected p clears .01, and the reported effect size is the margin, accuracy relative to the null mean rescaled to the interval up to 1, zeroed when the significance bar isn't met. The add-one permutation floor fixes the smallest achievable corrected p at $18/2501 \approx .0072$ for 410M and $14/2501 \approx .0056$ for 1B; those two constants generate all the gate arithmetic below.

The basis and the starving split. The basis is what distinguishes this instrument from a standard probe. For each capability, the specification names the tuple of surface components a lookup strategy would key on, analyzed against the tokenizer rather than against intuition about the task: Pythia's BPE splits numbers into digit chunks, so an operand-level basis under-specifies what the readout can key on for any digit-local label, and digit-local targets are banned at design time rather than starved badly (Experiment 2's addition lesson). Given the basis, a starving split draws a held-out value set per component. Validation items are those whose components are all held out; training items those whose components are all kept; items mixing the two are discarded from the fit entirely (Figure 1). A readout keyed on any component, or on an additive combination of per-component scores, therefore meets only unseen values at validation and scores chance by construction. Splits must hold out at least 15 values per component, produce at least 300 validation items, and keep every probe class present on both sides, enforced by seeded rejection with a bounded number of redraws; each of the five probe seeds redraws which values are held out, so a capability's five fits starve five different corners of its basis. Two variants cover geometries the plain construction handles badly: when components share one value space, a single holdout set drawn from the union applies to every position, closing the leak where a value held out at one position remains learnable at another; and when the label is a function of the basis value with rare classes, the holdout is drawn proportionally within label groups, which keeps rare classes on both sides. Capabilities that can't satisfy the constraints for all five seeds are ejected before the freeze with the failure recorded; five candidates died that way, one because its entire item space was smaller than the permutation count.

The gates. Four controls surround the scored fits, each with a tolerance derived from the floor arithmetic rather than set by feel. The untrained-weights gate runs the identical pipeline, splits and all, on a twin of each model with weights at seeded random initialization. Its expected clean-machinery fire count is the number of fits times the family-adjusted floor rate, about 1.6 fires across this campaign's 250 untrained fits, and the gate tests the observed count against that rate binomially instead of demanding zero. Fires structurally above the floor mark real leaks: the capability is dropped, and the experiment's own rules decide whether enough battery survives (the floor here was twenty of twenty-five; Section 4 records the outcome, and Section 6 what the closeout found wrong with the per-fire rule as frozen). The shuffled-label gate refits trained activations against labels permuted by a seeded generator under the same split geometry, permuting only the fit labels (the ordering as corrected mid-campaign; Section 4.2 discloses the original defect); it carries the same binomial tolerance, and a structurally-above fire there is grounds for declaring the pipeline unsound. The known-present gate points the instrument at capabilities that plainly exist, an entity-tracking task and a copy control, and requires seed-majority detection at both sizes with a mean starved margin of at least 0.2 at 1B: an instrument that can't see what's there has no business ruling on what isn't. An argmax gate requires the copy control to be generated correctly greedily at least 90% of the time at both probe sizes, certifying the generation-side harness the battery's inclusion decisions depend on.

Determinism. Probe fits were distributed across three machines with three BLAS stacks, and a box's results counted only after it reproduced a reference fixture through the real probe path bit for bit against




Figure 1: The basis-starved split. Items live on the grid of surface-component values. Probe training uses only items whose components are all kept; validation only those whose components are all held out; mixed items are discarded. A readout keyed on any component meets only unseen values at validation and scores chance by construction.

the Mac-computed reference, with no debugging tier in between; arm64 Accelerate, x86-64 OpenBLAS and Windows all matched under single-threaded BLAS, and the campaign’s 770 fits merge idempotently from any box that passed. The infrastructure is part of the claim: every number in Sections 4 and 5 recomputes from committed fit files. The fits themselves do depend on which machine produced them, by at most one to four validation items on 15 of the 39 remote fits, a smaller guarantee than the fixture gate appeared to give; Appendix C carries the re-fit and the rule it promotes, and no verdict or taxonomy assignment moves.

Table 1. The scored battery as frozen ($n = 25$), grouped by gate-1 fate: task, probe label, surface basis, and starving-split parameters. Held-out counts are over the values observed in the item pool and are means across the five probe seeds, as are the validation / training / discarded item counts. Three basis descriptions are abridged from the committed specification strings (the mapping is in `paper/table1_data.py`); the item files carry the verbatim text.

capability	task	probe label	surface basis	split (held-out / observed values)	val / train / disc.	fate
bin2dec	binary to decimal, 8-12 bits	value mod 10 (0-9)	bit-string (3968 significant 8-12-bit strings)	0.2 → 400/2,000	400 / 1,600 / 0	attrited
collatz2	two Collatz steps	ones digit of the FIRST-step result (0-9)	N token (~9990 values)	0.2 → 400/2,000	400 / 1,600 / 0	attrited
digitprod7	product of 4 digits (1-9), mod 7	the product mod 7 (0-6)	the number token ($9^4 =$ 6561 values)	0.2 → 400/2,000	400 / 1,600 / 0	attrited
isqrt	integer square root of N (100-9999)	ones digit of the root (0-9)	N token (~9900 values)	0.2 → 400/2,000	400 / 1,600 / 0	attrited
mod7_add	(a+b) mod 7, 2-digit operands	a mod 7 (0-6)	first operand token (90 values)	0.2 → 18/90	400 / 1,600 / 0	attrited
mod7_mul	(a*b) mod 7, 2-digit operands	a mod 7 (0-6)	first operand token (90 values)	0.2 → 18/90	396 / 1,604 / 0	attrited
mul3x1	3-digit x 1-digit multi- plication	tens digit of the product (0-9)	the 3-digit operand token (900 values)	0.2 → 167/835	392 / 1,608 / 0	attrited
numletter	(N x alphabet position) mod 26	the result (0-25)	N token (990 values); letters all seen	0.2 → 174/867	402 / 1,598 / 0	attrited
roman	roman numeral addition (values 1-99)	first numeral's value mod 10 (0-9)	first numeral string (99 values)	0.2 → 20/99	400 / 1,600 / 0	attrited
sq_mod7	($a^2 + b$) mod 7, 2-digit operands	a^2 mod 7 (quadratic residues: 0, 1, 2, 4)	the squared operand token (90 values)	0.2 → 18/90	403 / 1,597 / 0	attrited

capability	task	probe label	surface basis	split (held-out / observed values)	val / train / disc.	fate
units	metric conversion (power 1-3)	power of 10 (1-3)	unit pair (16 values; design's 3/16 holdout overrides the 15-value minimum)	0.2 $\rightarrow$ 3/16	378 / 1,622 / 0	attrited
unscramble	unscramble a 5-6 letter word	first letter of the solution	the solution word (multiset- unique pool)	0.2 $\rightarrow$ 150/736	407 / 1,593 / 0	attrited
weekday	day-of-week offset, N <= 499	N mod 7 (0-6)	offset token N (~499 values)	0.2 $\rightarrow$ 100/497	403 / 1,597 / 0	attrited
add3_mid	3-digit addition	middle digit of the sum (0-9)	tens-digit pair (100 values)	0.2 $\rightarrow$ 20/100	413 / 1,587 / 0	survivor
add_base8	octal addition, 2-digit operands	ones digit of the octal sum (0-7)	ones-digit pair in base 8 (64 values)	0.2 $\rightarrow$ 13/64	409 / 1,591 / 0	survivor
antonym	which option is the opposite of the cue	answer position (1-4)	the cue word (the cue-answer association is the lookup)	0.2 $\rightarrow$ 24/117	408 / 1,592 / 0	survivor
base7	write N in base 7	N mod 7 (0-6, the last digit)	N token (~9990 values)	0.2 $\rightarrow$ 400/2,000	400 / 1,600 / 0	survivor
caesar	decode a Caesar shift (k stated, 1-5)	first letter of the decoded word	(first cipher letter, shift) combo (~130 values)	0.2 $\rightarrow$ 44/110	795 / 1,205 / 0	survivor
clock24	24-hour clock, +D hours (D 25-499)	(H+D) mod 24 (0-23)	offset token D (475 values); mod-of-offset sibling of weekday	0.2 $\rightarrow$ 95/471	405 / 1,595 / 0	survivor
count_div7	count of multiples of 7 in [a, b]	the count (~4-18)	both endpoint tokens (shared value space)	0.45 $\rightarrow$ 446/882; 446/934	824 / 1,201 / 1,975	survivor
mod13	(a+b) mod 13, 3-digit operands	a mod 13 (0-12)	first operand token (900 values)	0.2 $\rightarrow$ 164/817	391 / 1,609 / 0	survivor

capability	task	probe label	surface basis	split (held-out / observed values)	val / train / disc.	fate
oct2dec	octal to decimal, 3-4 octal digits	value mod 10 (0-9)	octal string (4,032 values); value-mod- 10 sibling of bin2dec	0.2 → 400/2,000	400 / 1,600 / 0	survivor
odd_one_out	which of 4 words is not like the others	answer position (1-4)	all 4 words (shared components over the category vocab)	0.45 → 72/130; 72/149; 72/151; 72/136	338 / 741 / 6,921	survivor
reverse_string	reverse a random 4-6 letter string	last letter of the input (26)	final BPE chunk of the input (chunks, not strings, are the lookup unit)	0.2 → 130/650	408 / 1,592 / 0	survivor
sub3_mid	3-digit subtraction (non- negative)	middle digit of the difference (0-9)	tens-digit pair (100 values)	0.2 → 20/100	408 / 1,592 / 0	survivor

4 Two preregistered stress tests

The evidence for the prescription isn’t a benchmark built to flatter it. It’s two preregistered experiments that the controls killed, in sequence, for different reasons, before either one was allowed to touch its outcome variable. Both designs were frozen and tagged before any data existed; both adjudications ran as committed code against written rules; both verdicts were projected in a timestamped ledger before the formal report executed. The full record, including every number below, is in the tagged repository history.

Both serve the program described at the end of Section 2: probe results at 410M and 1B had to be committed before any larger model was queried, so the instrument had to be trustworthy on its own, with no outcome data available to rescue it. Neither experiment ever queried the larger models.

4.1 Experiment 2: standard splits, and what a random network decodes

The first experiment probed twelve capabilities on Pythia at 410M and 1B: modular arithmetic, two- and three-digit addition, two-digit multiplication, roman numeral conversion, weekday offsets, unit conversion, string reversal, unscramble, acronym formation, category counting, and a substitution cipher. Roughly two thousand probe items per capability; linear probes swept over (layer, token) candidates against a 2,500-draw permutation null with a Bonferroni family correction; five probe seeds per cell. The design included the control this paper argues for: the identical pipeline run on an untrained twin of each model, same architecture and tokenizer, weights at seeded random initialization. What it lacked was any defense in the split. Probe train and validation rows were drawn from the same item pool the standard way, stratified but otherwise random.

The untrained control fired on everything. All twelve capabilities, both sizes, every seed: 120 of 120 fits significant, each at the family-corrected permutation floor ($p = 18/2501 \approx .0072$, meaning the observed

accuracy beat all 2,500 label permutations). These weren't marginal detections. At 410M, seed-0 margins were 1.000 for unit conversion and mod-7 arithmetic, 0.992 for two-digit addition, 0.982 for roman numerals, 0.970 for weekday offsets; the weakest, the cipher, still cleared 0.205 against its own null. Replication at 1B fired 60 of 60 with near-identical margins, and the letter tasks got stronger (acronym 0.262 to 0.345, reversal 0.350 to 0.390): doubling the width hands the linear readout twice the random features, so the bigger probe-side model deepens the confound instead of escaping it.

The mechanism is what the reservoir-computing literature would predict. A fixed random network is a feature expansion; with roughly 2,000 examples against 1,024 features, a linear readout can assemble any deterministic function of the visible tokens by regression, and every capability label in the battery is a deterministic function of its prompt. A group-split diagnostic made this concrete for mod-7: the untrained probe's accuracy survives operand pairs it never saw (1.000) but collapses to chance on held-out operand values (0.129–0.142 against a 0.143 majority baseline). It isn't computing arithmetic; it's storing a per-operand-token offset table and mixing it additively.

The diagnostic's sharper finding is about the trained network. On the same held-out-operand split, the trained mod-7 probe collapses to 0.009, below chance, the signature of a lookup that's systematically wrong off its table. The trained probe's perfect in-distribution margin was lookup too. And a trained-minus-untrained gap score, the obvious repair, turned out to saturate exactly where it was needed most: five of twelve capabilities had both readouts at ceiling, because 1,500 training rows memorize a 90-value token table perfectly on either network. Two-digit addition explained the residue: Pythia's BPE splits numbers into digit chunks, the addition label is digit-local, and an operand-level holdout can't starve a digit-level basis, so even the random network generalized across operands (0.983). The lesson that became the successor's design rule: the split has to starve the basis the readout actually uses, and that basis is a per-capability, tokenizer-specific object.

The frozen attrition rule dropped all twelve capabilities; the battery floor was ten; the verdict was INSUFFICIENT_DATA with zero outcome-side queries. Two smaller findings from the same closeout carried forward. The shuffled-label control fired once in 120 fits, which the design treated as an abort condition; the closeout ruled it the test's own designed false-positive rate (expected fires ≈ 0.86), a zero-tolerance rule mistakenly placed on a nonzero-rate test. And one of two argmax positive controls failed at 410M (0.338 against a 0.9 bar) because its reliability had been assumed from another model rather than measured, which is the same assumption-transfer mistake in a different costume. Both reappear in Section 6.

4.2 Experiment 2b: basis-starved splits, and the class underneath

The successor was designed around the diagnosis. Every candidate capability now declares a surface basis: the tuple of prompt components a lookup strategy would key on, analyzed against the tokenizer (the operand token, the solution word, the unit pair, the letter-shift combination). A starving split holds out a set of values per component; validation items are those whose components are all held out, training items those whose components are all kept, and mixed items are discarded. A readout keyed on any basis component therefore scores chance on validation by construction. Capabilities with no feasible starving split, including anything digit-local, were excluded at design time. Gates got binomial tolerances derived from the permutation-floor arithmetic instead of zero-tolerance rules. The battery floor rose to twenty of a frozen twenty-five.

Before the freeze, the construction was rehearsed against Experiment 2's own activations, replaying two worlds known to be pure lookup through the new splits. Both went silent. The class that killed Experiment 2 is demonstrably closed by starving; that much is a positive result for the construction, and the campaign data repeats it at scale: the mod-7 family that fired at margin 1.000 under standard splits fired at seed-maximum margins of 0.10 to 0.26 under starved ones, scattered one-to-two-seed events at the weak edge of detection.

The campaign fit 770 probe cells over eight days on three machines under Section 3's determinism gate. One implementation bug surfaced mid-campaign: the shuffled-control stage permuted labels before building the split, which is undefined for label-stratified splits. The fix was argued from mechanism in the ledger before rerunning (the instrument's own permutation null already defined the correct order: split from true labels, permute only the fit), and the affected stage was refit uniformly. It's recorded here because hiding it would misrepresent what preregistered infrastructure work looks like.

The untrained control then found the class underneath the lookup class. Eighty-six of 250 untrained fits fired structurally, spanning thirteen of the twenty-five capabilities. These fires don't resemble the tolerated floor class: per-capability counts are binomially impossible against the floor rate (a capability's ten fits expect 0.064 fires; the strong leakers fired on ten of ten; Figure 3a), and the margins reproduce across the two independently initialized untrained models to within a few hundredths (the strongest, collatz steps, at 0.74–0.78 at 410M and 0.75–0.82 at 1B). Whatever carries these labels through a random network, it isn't the particular weights, and starving the basis didn't remove it. Section 5 dissects what it is.

The frozen rules did what they were written to do. Attrition removed all thirteen; twelve survivors sat below the floor of twenty; the verdict was again `INSUFFICIENT_DATA`, again with zero outcome-side queries. The other gates behaved: the known-present control confirmed the instrument sees capabilities that plainly exist (entity tracking at mean starved margins 0.330 and 0.281 across sizes; a copy control at 0.997), and the shuffled control's two fires in 250 sat at its designed rate (count-test $p = 0.538$). The argmax positive control repeated Experiment 2's assumption-transfer lesson in miniature: measured at 0.994 on the prior battery's items, it failed the 0.9 bar on this battery's items at 410M (0.868, upper confidence bound 0.896), a reliability that had been assumed to transfer and didn't. The verdict was projected in the ledger from gate-1 data three days before the frozen report ran; the report confirmed it without surprise. For a preregistered program, that's the intended shape of even a failure.

One honesty note on the adjudication itself: the frozen report contained two drafting defects, caught at closeout, and neither changed the verdict. Section 6 treats them as first-class results.

5 The leak taxonomy

Thirteen capabilities survived basis starving on a network that has never seen data. The question with reuse value is what property of a task predicts that. The detection facts below are preregistered results; the mechanism attributions are post-hoc hypotheses, formed by reading each capability's label definition against its basis, and they should be tested by construction in future batteries rather than trusted from this one.

The pattern, stated up front: a label leaks when it is a low-complexity function of surface fragments that the starving split cannot remove, because the fragments are shared across held-out basis values. The basis holds out values; it cannot hold out the alphabet the values are written in. Labels that are morphological or low-modulus functions of that alphabet ride through. Labels that require jointly composing the starved components do not.

The strong tier makes this concrete. Roman numeral addition leaked at margins 0.64–0.82: its probe label is the first numeral's value mod 10, and roman numerals write value-mod-10 directly into the numeral's suffix (I, II, III, IV...), so the label is literally printed on the surface in a shared sub-alphabet that value-level starving can't touch. Unscramble leaked at 0.18–0.22: its label is the solution's first letter, and that letter is one of the five or six letters sitting visibly in the prompt; a random expansion plus English letter statistics narrows the candidates. The Collatz task leaked strongest of all (0.74–0.82 across both sizes): its label is the ones digit of the first Collatz step, which is (up to the parity branch) a function of $N \bmod 20$, and $N \bmod 20$ is carried by the final digit tokens that all operand values share. Integer square root (0.36–0.44) is a magnitude-binning label: the ones digit of the root is constant on wide contiguous bands of N , and magnitude is legible from leading digits and token length. Digit-product mod 7 (0.19–0.30) is a fixed function of four visible digits that random nonlinear mixtures partially express. In each case the leak is not memorization of held-out values, which starving forbids, but computation over a shared surface vocabulary that starving cannot forbid.

The weak systematic tier says the same thing at lower amplitude: number-times-letter mod 26, three-by-one-digit multiplication's tens digit, binary-to-decimal mod 10, weekday mod 7, and the mod-7 family all fired at margins between 0.06 and 0.26 on scattered or systematic seeds. These labels need joint composition of digits that random features mostly can't provide, and starving removes the rest; what survives is the weak shadow of partial digit statistics. Unit conversion is the instructive oddball: it fired on exactly one probe seed at both sizes (0.39 and 0.52), and its basis has only sixteen values with three held out. Its label, the power of ten, is written in the unit's prefix morpheme (kilo-, centi-, milli-), so whether the leak shows depends on

whether a particular holdout draw splits a prefix family across the train-validation boundary. A leak can be a property of the draw, not just the task, when the basis is small.

The trained side of the same fits, computed from the closeout data (Figure 2), separates the thirteen into fates that matter for battery design. For six of them, training adds real structure on top of the leak: the trained starved margins exceed the untrained ones by up to 0.25 (integer square root by 0.21 and 0.25 at the two sizes, Collatz by 0.17, roman by 0.16, unscramble by 0.10, three-by-one multiplication and unit conversion by smaller and less stable gaps). These are tasks the models genuinely partially compute below threshold, sitting on bases that leak; with the label redefined off the surface carrier, they’re candidates for rescue. For digit-product and number-times-letter the gap is 0.03 to 0.06, and what the probe reads is mostly the surface floor with training adding little. And for five, the gap runs slightly negative: the mod-7 family, binary-to-decimal, and weekday show trained starved margins at or near zero while the random network fires weakly. Training appears to specialize features away from the generic surface mixing a random expansion provides, so on a starved split the trained network is less decodable than noise. All three regimes produce identical absolute-significance results under standard splits; they’re only distinguishable because both controls exist.

The twelve survivors sharpen the contrast from the other side. Their untrained starved margins are 0.000 in every cell, both sizes, all seeds. Eight are also silent or near-silent on trained weights at these scales (seven at exactly zero, one at a single 0.009), which is what the parent program would have scored against the eval ladder. Four show trained structure where the random network shows none: string reversal at 0.57 and 0.68 across the sizes, antonyms at 0.58 and 0.45, odd-one-out at 0.32 and 0.27, count-divisible-by-7 at 0.12 and 0.11. String reversal is the cleanest before-and-after in the record: under Experiment 2’s standard splits its untrained twin fired at 0.34–0.39, and under a basis starved on the solution word the untrained margin is exactly zero while the trained margin holds. That’s the construction doing precisely what it was designed to do, on a capability whose label isn’t written in a shared surface alphabet.

What the survivors share is that their labels sit at the far end of a composition: the middle digit of a three-digit sum with its carry structure, a base-7 or base-8 or octal conversion, a Caesar rotation, a modulus of 13 rather than of the digit alphabet. The design guidance this yields is checkable rather than aesthetic. Before including a capability, write down its label as a function of the prompt and ask what the cheapest surface approximation of that function is: whether the label appears in the prompt’s visible fragments, whether it’s constant on magnitude bands, whether it’s a small-modulus function of single tokens, whether the basis vocabulary is small enough for one holdout draw to matter. Then measure it: run the untrained twin on the starved split at generation time, before inclusion, and let the acceptance test rather than the designer’s intuition say whether the basis is starved. Every mechanism above was legible in the label definitions before the campaign ran. I found them afterward because I hadn’t yet learned to ask the question in that order. Section 7 shows what asking in that order looks like.

6 Calibration lessons for frozen criteria

Preregistration moves analysis decisions ahead of the data, but the decisions still get written down as thresholds and code, and both can misimplement the design they encode. A defect frozen in good faith binds like any other frozen choice. This section catalogues the ones this program found in its own frozen instruments, each with the rule that would have caught it before the freeze. One fact holds for every ruling below and is stated once: the mechanism argument behind it was written before, or demonstrably independent of, the outcome it touched, and every projection was timestamped in the ledger before the frozen report ran. That is what makes a ruling adjudication rather than rescue. Lessons 1 to 4 come from the campaigns this paper reports; lessons 5 to 11 come from the successor experiments in the supporting record, each named by its tags.

The first lesson: a signature bar is derived from the mechanism that generates the events it classifies, never from an SD intuition. Experiment 2b tolerated occasional control fires at the permutation floor as the pipeline’s designed false-positive rate, and its frozen report classified a fire as tolerable when it sat at the add-one floor and within 3 SD of the null mean. But a fire at the floor beat all 2,500 permuted fits, and the expected maximum of 2,500 null draws lies about 3.5 SD above the null




Figure 2: Trained versus untrained starved-probe margins per capability (means of 5 seeds), both model sizes; attracted capabilities above the divider, survivors below. An orange ring behind a blue dot marks coincident values. Survivors' untrained margins are exactly zero in every cell; four survivors show trained structure the random network lacks entirely; and the top of the attracted block shows training adding structure on top of a leak, the rescuability signature described in the text.

mean (Figure 3b). The tolerated class is nearly empty by construction, and a clean null was expected to hand the predicate about 1.6 misclassified fires per 250, at least one with probability about .8. The shuffled control then behaved exactly as designed, two fires in 250 with a count-test p of .538, at 3.6 and 4.7 null SD,




Figure 3: Gate arithmetic. (a) Untrained structural fires per capability against the floor-rate expectation of 0.064 fires per capability; the shuffled gate’s 2 fires in 250 sit at its designed rate. (b) The floor-signature contradiction: a fire at the permutation floor beat all 2,500 permuted fits, so it lives where the maximum of 2,500 null draws lives (mean ≈ 3.5 SD), while the predicate tolerated only fires within 3 SD of the null mean, a region the maximum reaches with probability $\approx .03$. The two observed shuffled fires, at 3.6 and 4.7 SD, were both misread as structural.

and the predicate read both as structural and projected a pipeline abort on a control whose own count test was passing. The closeout ruled the fires floor-rate events and the predicate a design-level miscalibration. Experiment 1 had already paid for one instance of the same class, the misspecified magnitude criterion of Section 3’s provenance note, so the rule is standing rather than advisory.

The second lesson: adjudication code freezes together with fixture tests derived from the design document’s own worked examples, plus one synthetic case per preregistered provision. The design gave the untrained gate exactly one remedy, attrition of the leaking capability, and reserved pipeline abort for the shuffled gate. The report’s code instead pooled every structural fire into one binomial count test against the floor rate, which trips at seven fires in 250, while a single leaking capability contributes ten fits: any attrition event was an abort, and the provision could never have operated for any possible data. On the observed 86 fires the pooled test returned $p = 6.5\text{e-}117$ and declared the pipeline unsound; the closeout ruled the deviation against the design text and closed the gate as attrition times thirteen, safe to rule because both readings end in the same place, twelve survivors against a floor of twenty. No test had executed the provision before real data did. “One leaking capability” must yield attrition-without-abort before the freeze; both of 2b’s implementation deviations, this one and the shuffle-before-split bug of Section 4.2, fail that fixture immediately.

The third lesson: gates whose inputs are committed before the freeze get adjudicated before the freeze. The argmax control’s 410M failure was computable from the committed inclusion record eight days before the campaign spent its compute; the frozen report surfaced it at campaign end because nothing required checking a gate the moment its inputs existed. The skipped check cost nothing only because the battery died of something else first.

The fourth lesson: the number reported beside a corrected test is a selected value, and selection is not what the correction protects against. The instrument sweeps fourteen to eighteen candidate sites, tests each against a permutation null, applies a Bonferroni correction across the family, and reports the site with the smallest corrected p , ties broken by accuracy. The correction controls the false-positive rate of the detection. The accuracy printed beside it is the accuracy at the site the selection chose, and nothing in the procedure makes it an unbiased estimate of what the representation contains. The records stored only

that one number per fit, so the campaign could not quantify its own exposure; re-fitting every candidate from the stored activations recovers what was discarded (of 770 fits, 755 reproduce their stored accuracy exactly at the selected site, and the measurement uses those). The untrained stage is the cleanest read, because every candidate there is measuring the same nothing:

stage	selected site	mean over candidates	inflation
untrained	0.2307	0.1808	+0.0499
shuffled labels	0.1568	0.1282	+0.0287
scored battery	0.3046	0.2215	+0.0831

On a randomly initialized network the reported site reads 1.28 times the average site, positive in **246 of 247 untrained fits**; the exception is a fit in which all fourteen candidates returned the identical accuracy. The selected site was also the highest-accuracy site in 92 to 100 percent of fits, so in practice the rule reports a maximum whatever its tie-breaking says. The successor battery, replayed the same way, gives 0.1393 against 0.1045, a ratio of 1.33, positive in **227 of 227** fits. The artifact was not holding that battery’s result up: its primary rank correlation, recomputed with the profile mean in place of the selected site over the same 22 rungs, moves from .431 to .456 (neither is the frozen verdict, which normalizes against a permutation null; the contrast is the point). That is the opposite of Section 8’s floor defect, where the correction cut the correlation from .368 to .200. A selection artifact is not automatically load-bearing; what it changes is what the number means. The prescription is not to prefer the mean, which dilutes a signal localized to one block, but to **report the per-site profile, with the untrained twin’s value beside each site**. Every candidate was already fit in order to take the maximum, so the cost is storage, which is what we failed to spend.

One corroboration from outside this paper’s evidence base, with its weakness attached: the same two-gate apparatus, applied per site, ran on a synthetic formal language with a percolation threshold, on forty models trained below it where the capability cannot form, and both controls were silent, zero fires in 480 untrained and 320 trained sites. That bounds the worry that an apparatus which always finds a leak is measuring its own sensitivity, but weakly: that probe’s split is entity-wise by construction, the starved condition of Section 4.2 for free, and finding a hygienic system hygienic is close to circular. Section 7’s clean gate on this paper’s own battery is the stronger evidence.

The freeze as an adversarial instrument. Everything above was caught by gates running against data. The successor experiments added a place to catch defects before any data exists: a three-session protocol, design, build and freeze, separated by context boundaries, where the freeze opens with a cold re-read of design and build under one instruction, assume the tree still contains a class defect and find it. The assumption has paid out every time. One freeze found a mass statistic crediting set-level lexical priming, a position-blind primer with nothing of the capability in it scoring 500 of 500 ($p \approx 3e-151$), and the campaign’s own data later showed the superseded form would have decided the verdict; the next found the leak gate re-rendering its prompts from live item files with no hash pin, and a byte-continuity attestation never checked against the tree the analysis pooled (`exp3-*`, `exp3c-*`). The pattern is the same each time: the defect lives on a path the test suite substitutes for, so a fixture suite at a 100 percent mutation kill rate can still miss it. “Find the class defect” is a search of the seams between the tested and the real, and it belongs on the checklist. Lessons 5 to 11 are what those freezes and the campaigns after them found, each stated as its rule with the evidence cut to what carries it.

The fifth lesson: a frozen criterion must be total over the alphabet it scores. The exact-match verifier behind every generation number in this program normalizes the model’s text and indexes the first whitespace token of the first line, which assumes a token exists. On a first line that is bare whitespace wrapped in punctuation it raises instead of returning. The defect ran 4.2 million temperature-1 draws across four closed experiments without firing, then stopped a campaign on the one draw in 384,000 that landed in the class. It could only raise, never mis-score, so every closed verdict stood; the fix, a measurement-side wrapper mapping the crash class to a non-fire, was proven verdict-preserving by recomputing the full committed corpus to identical counts before the campaign resumed. Adversarial review interrogates what a criterion credits; totality is a different property and needs a different test. Property-fuzz the frozen verifier over the

emission alphabet at the freeze and require a verdict, not an exception, on all of it. (Record under the `exp3c-*` tags.)

The sixth lesson: a preregistered power table is a claim about the alternative’s shape, not only its size. The rank-prediction experiment declared itself underpowered in advance, computed power 0.26 against a house bar of 0.75, ran anyway, and rejected its null at $p = 1.6e-4$. A p-value that far past a test’s stated reach is evidence the alternative was specified wrongly, and the error is recoverable from the committed power record: the frozen alternative spread rate mass over the thirteen items observed to succeed, while the structure the data carried was a forty-five-item equivalence class of which those thirteen were an unrepresentative sample, a target roughly six times larger than modelled. The sensitivity analysis varied how strongly an item-level story was diluted and never asked whether the story was item-level at all; the freeze had printed the predictor’s tie structure, unread. When a frozen predictor’s ties collapse items into classes, compute power over the realized classes as well as per item, and treat a large discrepancy as the design question it is. The corollary cuts against intuition: a coarse instrument is not a weak one. Tie structure bounds what a statistic can express, not how large an effect it can detect. (Record under the `exp3d-*` tags.)

The seventh lesson: a conservative rule is conservative relative to a test, and “applied identically” can flip its sign. The program’s leak rule voids any draw that emits a gold answer occurring verbatim in the rendered prompt. For a fire count the rule can only lower the number, so it is safe wherever a higher count argues for the claim. The shortcut experiment applied it, in its own words “identically”, to competitor strings inside a relative test that designates one of an item’s one-edit neighbours “the reverse” at random, conditional on the observed counts. Zeroing a void competitor’s count there removes a slot that could have scored and leaves the reverse’s share larger: on a three-item instance the designation p is $1/4$ with the competitor live, $1/8$ with the item excluded, $1/24$ with the competitor zeroed. The closure was to drop any item carrying a void target from the relative test; on the committed battery no target was void, so no number moved. “Apply the same rule” describes an operation, not its effect, and the effect depends on whether the statistic is an absolute count or a share. Classify every exclusion rule by the test it feeds and re-derive its direction there. (Record under the `exp3e-*` tags.)

The eighth lesson: a frozen verdict must be able to deliver its own refusal terminal from the tree the runner actually leaves behind. The sampling-ladder experiment’s verdict tree opened with `INSUFFICIENT_DATA`, fired when the byte comparison against a committed stream found any difference. The runner implemented the halt correctly, parking the differing rows and withholding the rung’s normal draws file so a resumed campaign could never mistake a halted rung for a finished one. The frozen analyzer loaded that tier before it read the gate, and its loader refused an incomplete tier with a hard error: on exactly the tree a halted campaign leaves, the verdict raised `FileNotFoundError` instead of saying `INSUFFICIENT_DATA`. Every synthetic world had been built from the analyzer’s own writers and none from the runner’s failure path. The closure reads the gate records first and returns the terminal before any tier loads; the defect never fired on the real campaign, which reproduced all 128,000 committed draws byte for byte. A verdict that crashes where it should refuse has an unreachable terminal, and an unreachable terminal is a preregistered outcome the instrument cannot honour. For every refusal terminal, build the full-shape world with the producer’s own failure path, not the analyzer’s writers. (Record under the `exp2d-*` tags.)

The ninth lesson: the target is part of the instrument, and a probe’s silence is a statement about a label and a split until the same activations have been read against a label the representation could carry. Section 7’s battery left two rungs the probe read as silent at both sizes and the outcome read as rising by 2.8B: three-digit subtraction scored on the middle digit, and the next term of an arithmetic sequence. Under the program’s own discriminator that is the interesting pattern, and the projection had named it as the sharpest disconfirmation available. Laying the records side by side showed the instruments had never read the same quantity: the sequence rung’s probe label was the next term’s residue modulo 7, chosen so that no lookup table could score it and validated on held-out basis values so that a memorised basis could not either; the generators had been scored on the full answer; and the probe’s committed best site for both rungs was a layer-0 constant feature at corrected $p = 1.0$. Experiment 2f put the three instruments on one label per rung, the middle digit and the last digit, with one floor and one bar, on the same weights and the same 500 items. The probe read the digit in every cell, .23 to .27 against bars of .17 to .19 with the untrained twins at .04 to .10, while the generators read it in one cell of four; and the activations that

read the last digit at .262 and .270 read the mod-7 residue at .124 against a .212 bar. Nothing in the probe machinery was wrong. A label chosen to defeat lookup and a split chosen to test generalisation, neither asked to justify itself as a measurement of presence, together made a rung look like a percolation case: the kinship with Section 8 is exact and on the other side of the correlation, an outcome variable that credited format as capability there, a predictor variable that scored generalisation as presence here. Before a silence is read as absence, read the same activations against the most direct label the task admits, and if the two disagree the silence belongs to the target. A battery that starves every label without a matched-label control has built a generalisation test and called it a presence test. (Record under the `exp2f-*` tags.)

The tenth lesson: a fixture suite certifies an instrument’s logic, and the seam between fake and real inputs is exactly where it stops; every execution path a frozen instrument will take on campaign day must run once, end to end, against real inputs before the tag. The cross-family experiment carried the largest freeze battery in this record, 284 tests, 118 of 118 mutants killed, a hundred thousand fuzzed inputs through the scoring criterion, and its model-contact rehearsals passed. One path had never touched a real model before the tag, the sweep’s checkpoint loader, because loaders are model contact and the suite drives them through injected fakes. The campaign’s first model load crashed it: the installed library reads a generation config from the model directory after the weights, falls back to `config.json`, and special-cases that one filename in a local directory to return None where any other absence raises, so a directory holding only weight shards loaded 355 of 355 tensors and died four frames deep in someone else’s fallback. The predecessor’s loader had written `config.json`; the generalisation that dropped the write was made for a good reason and reviewed by every instrument that could not see it. Nothing scored; the preregistration tag was re-cut at the fix commit, a disclosure the public record carries. This is the third appearance of one class, after the fifth and eighth lessons, and the class has a name: the production path was exactly the untested path. Before the tag, enumerate the paths campaign day will take, every loader, producer and halt, and require of each one real end-to-end execution, however small. (Record under the `exp2i-*` tags.)

The eleventh lesson: an analyzer’s inputs include the code the interpreter executes on its behalf, and a sweep of what the analyzer reads cannot see what it imports. The mechanism experiment was analysis-only, every number a re-read of committed bytes, and its freeze inherited the full pinning discipline: twenty-six named modules hashed against literals, a manifest of 2,621 data files, the instrument files bound to the tag by blob, and a read sweep over every file-opening primitive reporting zero unpinned reads. A scan of the interpreter’s module table after one real-tree run found twenty-three files under the experiments tree loaded into the analyzer’s process and covered by no pin: every package `__init__.py` on the import chain, the predecessor battery’s item generators reached through the pinned item loader, and seven further helpers. The sweep could not have seen them by construction, because the import machinery reads a module’s bytes before any wrapper is installed. Two lines in an empty `__init__.py` rebinding one functional moved the primary statistic from .1311 to .1330 with every gate passing and the sweep still reading zero. Closed with a twenty-eight-file pin over the resolved module table, checked at the analyzer’s entry and again at its exit. This is the fourth generation of one lineage, every verdict input pinned at analysis time, each generation pinning one surface over from the last. Pin the set of repository modules the interpreter has loaded, not the set the analyzer names; a pin on named modules certifies what the author remembered to import. (Record under the `exp2j-*` tags.)

7 The screen at inclusion time

Both campaigns above ran the untrained control after their batteries were built, and Section 4’s verdicts are the price of that ordering. The successor experiment ran it where P1 says it belongs: its battery was constructed under the screen and frozen in August 2026 (tag `exp2c-preregistered`), and everything in this section is probe-side screening data committed at or before that freeze. What the record shows is what the screen catches at design time, what the catches cost, and what a battery that passes looks like; Section 8 reports the outcome.

The screen is tiered so it can live inside the design loop. Tier 1 runs at candidate-design time: two untrained seeds at both sizes against 500 permutations, roughly an hour of compute per candidate, with the rejection

bar derived from the order statistics of the maximum of 500 null draws (Section 6’s calibration rule, applied prospectively rather than learned again). Tier 2 runs on tier-1 survivors before the freeze at the identical frozen configuration, five seeds by both sizes by 2,500 permutations, and the freeze declares those fits to be the campaign’s untrained-gate fits, so the gate that ended both campaigns in Section 4 is adjudicated before the battery locks rather than after it has spent its compute. Experiment 2b bought its taxonomy with a 770-fit, eight-day campaign; each discovery below cost a screening pass.

The first catch was the screen out-analyzing its designer. A base-12 conversion rung entered with the label “N mod 12, the last base-12 digit” and a written dumbest-baseline argument that a digit-suffix lookup fails structurally because 12 never divides a power of ten. True, and incomplete: N mod 12 decomposes by the Chinese remainder theorem into N mod 3, the decimal digit sum, and N mod 4, the last two decimal digits, both classic surface carriers. The tier-1 screen fired structural-abort on four of four fits. The replacement asks the identical question with the label moved to the digit sum of the base-12 representation, mod 5; 5 is coprime to 3, 4, and 11, so none of the surface-legible congruences reaches the new label, and the screen went silent on the same prompts and the same basis. That fire-to-silence contrast is Section 5’s closing advice operating in the intended order: the acceptance test, not the design argument, ruled on the basis, and it caught a carrier the argument never considered. The layered review then earned its keep on the replacement itself, catching a false coprimality claim (coprime “to 12, 10, and 11”; $\gcd(5, 10) = 5$) before it froze — two catchers, two different catches.

Three rescues then ran the constructive test Section 5 proposed. Roman numeral addition, Collatz stepping, and integer square root re-entered with labels moved off their named carriers: the sum of the two numeral values mod 7 in place of any numeral’s suffix, the second Collatz step in place of the N-mod-20-legible first, the gap between N and the largest square below it in place of the magnitude-banded root digit. All three passed the tier-1 screen that had caught their parents at untrained margins up to 0.82. A correct attribution predicts exactly this: name the carrier, move the label off it, and the untrained readout goes silent on the same task shape. It did, three times. Those three rows of Section 5 are upgraded from hypothesis to constructively tested; the remaining attributions stay hypotheses.

The screen also found a class Section 5 doesn’t have. Two rungs asked for the letter of a printed 8-letter string at a computed interior position, $p = ((i + j) \bmod 6) + 2$ and a product variant, with i and j printed alongside. The basis was the string alone; the position arithmetic cannot be starved, since i and j are printed and the map is low-complexity, and the specs disclosed the gamble in writing: the untrained screen would arbitrate whether a variable-index gather is decodable from random features. It is. Both rungs fired structural-abort on four of four fits, whether the position distribution was concentrated or near-uniform, and the family was ejected whole. The datum is genuinely new: string reversal’s fixed-position label reads exactly zero from untrained weights (Section 5), but a label at a surface-computable variable position leaks. Positional gather joins the taxonomy, at the cost of two screening passes instead of a campaign.

Not every catch came from the screen, and not every named carrier leaked. A six-word odd-one-out rung’s blessed six-component split satisfied its feasibility floors only degenerately: each seed cleared them only when the holdout swallowed one complete word category, making starved validation effectively leave-one-category-out, and the rung was re-blessed to a one-component basis before screening (minimum starved validation rose from 305 items to 2,336). The general lesson: a split can meet its committed counts while the draws that meet them change what the margin measures, so the holdout’s actual contents get inspected, not just its arithmetic. In the other direction, a six-option antonym rung came out of adversarial review with a named and measured carrier, wordlist slot-asymmetry giving cheap heuristics (answer length, edit distance to the cue, shared letters) about 0.20 accuracy against 0.167 chance, ledgered before its screen ran, with the screen preregistered as the adjudicator. The screen never fired; corrected p was 1.0 in every fit. A sequence-extrapolation family carrying disclosed on-surface mod-7 functionals of its printed terms was silent too. A measured 20-percent heuristic dose sitting on the silent side of the line is a boundary datum the taxonomy needs: the screen separates the carriers random features express from the ones they don’t, and “surface-computable” is not one thing.

One rung fell to the other inclusion bar. Candidates must be behaviorally absent at the probe scales, normalized argmax margin at 1B with a 95% upper confidence bound under 0.25, the bar frozen in 2b. A

Hamming-match-count rung read 0.282 trained accuracy at 1B, upper bound 0.324, over the bar; the plausible mechanism is modal-answer guessing over its mid-range count distribution. The frozen rule is mechanical and carries no mechanism escape, so the rung ejected, the mechanism went into the ejection record, and the power gate was re-run on the reduced battery before the freeze (0.769 against the 0.75 floor, still passing). When a mechanical bar and a mechanism story disagree, the bar governs and the story is recorded; the alternative is negotiating with frozen criteria one sympathetic case at a time. The freeze review shed a rule of the screen's own by the same principle: a tier-1 margin bar that could never fire (at 500 permutations the significance floor sits above the frozen alpha) was dropped so the frozen record would not imply a gate that cannot fire.

What came out the other side is a battery of 34 rungs in 16 families with its untrained gate adjudicated at the freeze: 220 frozen-configuration fits over the 22 new rungs, two fires, both from a single rung at the add-one floor with z inside the tolerated band, against a clean-machinery expectation of 1.41 (binomial $p = .41$); the twelve carried survivors stay all-zero from 2b. The argmax control was measured on this battery's own items, per the Section 4.2 lesson, and read 0.960 and 0.980 against its 0.9 bar. Set against Section 4's totals, 120 fires in 120 fits and then 86 in 250, this is the program's first battery to reach its freeze with the untrained gate clean, and Section 9's worry that an acceptance test nothing has passed may be unpassable now has the only answer it could get: a battery passed, by being built to. Whether a battery that passes the screen measures anything worth measuring is that experiment's question, and Section 8 reports what it cost us to get there.

8 The same error on the outcome side

The battery that passed the screen went on to be scored against an eval side: argmax accuracy on held-out items at 2.8b, 6.9b and 12b, with each capability's score normalized against a floor, exactly as the probe margins are normalized against a permutation null. The floor we chose was the one this paper argues for everywhere else, the untrained network. Randomly initialized weights, same architecture, same items, same prompt. It is the control we spent seven sections recommending, and on the outcome side it is the wrong one.

An untrained language model does not produce wrong answers. It produces malformed ones. Asked for a number it emits fragments, and the normalizer that extracts an answer finds nothing to score. So the untrained floor came back at or indistinguishable from zero on all thirty-four capabilities, for a reason that has nothing to do with any capability: the model cannot participate in the task's format at all. A trained model that has acquired only the format, one that has learned to emit a small integer after **A**: and nothing more, sits at one over the size of the answer space. Against a floor of zero, that entire guessing rate is scored as capability.

The consequence was not marginal. Of thirty-four capabilities, ten landed within noise of their own chance rate and were credited the whole of it. Modular arithmetic at modulus 13 scored .046 against a chance rate of $1/13 = .077$, below chance, and was credited a normalized margin of .071. The effect is sharpest where chance is highest, which is exactly where a reader's intuition is weakest: for multiple-choice items the chance rate is one over the number of options presented, and a four-option capability scored .208 against its .250 chance, below chance again, and ranked fifth of thirty-four on the outcome variable. Two capabilities out of thirty-four cleared their chance rate by a factor of two or more.

The generalization is the same one this paper makes about probes, one level along. An untrained-weights control answers *is this signal in the weights?* It does not answer *is this behaviour competent?* Those are different questions with different floors, and the second one's floor is a property of the task's answer space, not of the network. Substituting one for the other credits format acquisition as capability in precisely the way that standard-split probing credits surface statistics as structure. We wrote seven sections on the probe-side version of this error and then committed the outcome-side version in our own frozen operationalization, which is the strongest evidence we can offer that the control is not obvious once the question changes.

Two properties of the error determine how much damage it does.

It is not conservative. Re-scoring against within-answer-space chance rates, as a disclosed descriptive that could not touch the frozen verdict, moved the primary rank correlation *down*, from .368 to .200, and cut the count of capabilities with nonzero outcome scores from twenty-five to thirteen. The contamination was

manufacturing correlation, not masking it. A reader’s natural assumption, that a too-generous floor adds noise and biases toward the null, is wrong here, because the inflation is systematic rather than random: it scales with the inverse of the answer space, and answer-space size is correlated with task type across any battery organized by task family.

And it is invisible to every control this paper recommends. The untrained gate was clean. The shuffled-label gate was clean. The known-present positive controls passed at both scales. Nothing in the probing hygiene apparatus we describe detects an outcome variable whose floor is measuring the wrong thing, because the apparatus is pointed at the representation and the defect is in the behaviour.

The prescription is one line: **report both floors**. Any margin normalized against an untrained network should carry its within-answer-space chance rate alongside, and any claim that a capability is present above floor should say which floor. For multiple-choice items the option count belongs in the committed record, not recoverable by re-parsing question text after the fact, which is how we recovered it.

A third observation from the same scoring is worth stating because it runs the other way and is, if anything, the more useful result. The two capabilities carrying the second and third highest starved probe margins in the battery, character-level string reversal at two lengths, scored an outcome of zero at every eval scale, and as first reported that contrast was confounded twice: the probe decoded one character where the eval demanded all seven exactly, and the margins lived at 410m and 1b while the zeros lived at 2.8b and above. A follow-up experiment (preregistered, frozen, run once; tags `exp3b-preregistered`, `exp3b-closed` in the supporting record) removed both confounds at once: it scored the first character of the greedy continuation against the probe’s own label, in the probe’s own 26-way space, on the probe’s own models. On the same weights where the starved probe reads the answer’s first character at margins .6263 and .7725 (seven-character reversal, 410m/1b) and .5731 and .6749 (variable-length), first-character emission is .0520/.0280 and .0320/.0260 against marginal floors of .056 and .054, at or below floor in all four cells, while the same instrument scores a copy control at .9940 on the same weights. The eval-scale zeros stand as committed, including the one item of five hundred at 6.9b, which the follow-up reproduced byte for byte. The information is linearly decodable from the representation, survives basis starving, and clears the untrained gate, and the model does not emit even its first character above chance under greedy decoding — with the resolution stated: the design cannot distinguish floor from a true rate below .092, so this is no emission at that resolution, not a certified zero. Decodable is not learned, and on this evidence decodable is not *generable* either, now on the same weights, the same prompt, and the same label space. A probe margin licenses a claim about what a representation contains, not about what the model will do, and here the gap between those is the difference between the strongest signal in the battery and behaviour indistinguishable from chance.

9 Limitations

Everything here was demonstrated on one model family, one tokenizer, and two scales. The basis analysis is tokenizer-specific on purpose: what a lookup can key on is a fact about the token inventory, so no battery’s basis definitions transfer to another tokenizer without redoing the analysis, and the taxonomy’s specific mechanisms (digit tokens carrying parity, prefix morphemes carrying powers of ten) are Pythia-BPE facts first. The probes are linear. That choice is what makes the capacity arithmetic clean and the permutation null cheap, and the free-computation problem only widens for more expressive probe families, but none of the tolerances here have been worked out for them. The candidate family is fixed and narrow, two token positions and every third layer, so a silence is always relative to the sweep: a capability decodable only at an unswept layer or position reads as absent. Fires are absolute; silences are family-relative. The known-present gate bounds the worst version of this, since an instrument blind to what plainly exists fails it, but the asymmetry stands. The scales are 410M and 1B, and the one scale comparison in the record points the wrong way for comfort: doubling the width roughly doubled the random features and strengthened the untrained readout on the letter tasks, so the confound should be assumed to grow with model size, not wash out.

The untrained control is necessary, not sufficient. The program found its two confound classes sequentially, the lookup class and then, after starving closed it, the surface-statistics class underneath, and the honest induction from that sequence is that more classes exist. A control against a random substrate bounds what comes for free; it cannot flag a shortcut that random features can’t express but training finds, and on a

starved split such a shortcut reads as learned structure. Starving is the only defense in this design against that, and starving covers exactly the declared basis, nothing else. The taxonomy’s mechanism attributions are post-hoc hypotheses and are labeled as such in Section 5; three have since been confirmed by the constructive test proposed there (Section 7), and the rest remain untested.

The prescriptions generalize from a single program’s history, and the objection writes itself: these are lessons from one group’s mistakes on one model family, and a checklist induced from $n = 1$ is folklore with tags. Two things blunt it. First, the checklist’s claims are existence claims, and existence is what a case study can establish. Each item asserts that a specific failure mode is reachable and cheap, not how often the field reaches it, and every item carries the commit where it was reached. Second, the conditions here were stacked against these failures, not for them: every mistake in this record was made under preregistration, frozen criteria, and controls run on purpose, by a program conspicuously trying to avoid exactly this class of error. A failure mode that fires under that much deliberate care is not plausibly rarer in work that reports untrained baselines as decoration, claims zeros without confidence bounds, and selects sites without selection statistics. What $n = 1$ cannot supply is a prevalence estimate, and none is claimed.

Finally, the method’s positive record is one battery frozen and one campaign completed, not one hypothesis supported. Both campaigns of Section 4 ended INSUFFICIENT_DATA. The successor battery of Section 7 passed the screen, reached its outcome measurement, and returned a FAIL its own preregistered power analysis had already declared uninterpretable: 22 of its 34 capabilities returned zero probe margin and 9 of its 16 family blocks were entirely inert, leaving a rank test with seven effective blocks where its frozen power table had assumed sixteen. So the controls detect bad batteries, and a battery can be built to pass them, and the one that did went quiet rather than dirty. Surface computability of the kind Section 5 maps stayed dense in practice (the successor’s construction still ejected four candidates across three distinct mechanisms), and whether enough task diversity survives the screen *while retaining measurable spread* is now the open question, sharper than before and not favourable: this paper’s prescriptions buy a clean battery, and a clean battery is not automatically an informative one. A successor should gate on predictor spread before spending an outcome campaign, which is computable at freeze and was computable here.

10 Recommendations

The checklist below compresses Sections 3 through 8 into the form I wish the program had started with. It’s written to be adopted verbatim. Every item is design-time analysis or a reuse of machinery a probing study already has; the one new cost, fitting the untrained twin, is a second pass of a pipeline that already exists. Each item ends with the section that supplies its evidence.

Table T2. The probing-hygiene checklist.

1. Write every probe label as an explicit function of the prompt, and record what a lookup table and a random network should score on it before the capability enters the battery. (§3, §5)
2. Name the surface basis a lookup would key on, analyzed against the tokenizer’s actual token inventory, not against task intuition. (§3)
3. Reject label targets that are local to single surface tokens: a digit-local label under digit BPE can’t be starved. (§3, §4.1)
4. Starve the basis: hold out values per component, validate only on items whose components are all held out, discard mixed items, commit the feasibility counts, and eject capabilities that can’t satisfy them. (§3)
5. Run the untrained twin through the identical pipeline (prompts, splits, candidate sweep, permutation null, correction) as a pre-inclusion acceptance test, not a post-hoc baseline. (§4, §7)
6. Derive every control’s tolerance from the pipeline’s own false-positive arithmetic; no zero-tolerance rules on nonzero-rate tests; test fire counts binomially, per capability. (§3, §6)
7. Derive per-fire signature bars from the order statistics of the permutation maximum, never from an SD intuition. (§6)
8. Measure control reliability on the current battery’s items; never transfer a reliability measurement across battery versions. (§4.2)

-
9. Adjudicate at freeze time every gate whose inputs exist at freeze time. (§6)
 10. Freeze adjudication code together with fixture tests built from the design document’s worked examples, plus one synthetic case per preregistered provision. (§6)
 11. Report every zero as a Clopper–Pearson bound. (§4.2)
 12. Project the verdict in a timestamped ledger before the frozen report runs. (§4.2, §6)
 13. If fits are distributed, admit a machine only after it reproduces a reference fixture through the real probe path bit for bit, and treat that as necessary, not sufficient: certify a worker on a random sample of the units it actually produced, and put load-bearing fits on one machine. (§3)
 14. Inspect what the starving holdout actually draws, not only its counts; a split can satisfy its feasibility floors degenerately, which changes what the margin measures. (§7)
 15. Normalize a behavioural outcome against the task’s within-answer-space chance rate as well as against the untrained network, and report both; an untrained floor sits at zero on a generation task and credits format acquisition as capability. Commit the option count for every multiple-choice item at freeze. (§8)
 16. Gate on predictor spread before spending an outcome campaign; the power of a mostly-tied predictor is computable at freeze from the realized predictor alone. (§8)
 17. Do not claim a probe margin licenses a behavioural prediction. (§8)
 18. Store and report the per-site profile, not only the selected site, with the untrained twin’s value beside each site; every candidate was already fit in order to select one, so the cost is storage. (§6)
 19. Freeze a producer for every input the adjudication function takes, not only the ones carrying the primary statistic: every parameter has either a frozen producer or a fixture pinning its value, or the verdict is not frozen. (§6)
 20. Open the freeze adversarially: a cold re-read of design and build under the assignment “find the class defect,” aimed at the seams the fixture suite substitutes for. (§6)
 21. Property-fuzz every frozen scoring criterion for totality over the emission alphabet before freezing: on garbage it returns a verdict, never raises. Guard the measurement side only; a referent that crashes the criterion is a broken battery and stays a hard error. (§6)
 22. For every refusal terminal of a frozen verdict tree, build the full-shape world with the producer’s own failure path and require the analyzer to return that terminal from the tree as left. (§6)
 23. Before a frozen probe label’s silence is read as absence, pair it with the most direct label the task admits, on the same bytes, and attribute a silence the direct label dissolves to the target, not the representation. (§6)
 24. Enumerate the execution paths a campaign will take, every loader, producer and halt, and run each once, end to end, against real inputs before the preregistration tag. (§6)
 25. Write pre-campaign checks against invariants, not snapshots, and re-run every cold verification tool after each stage lands, not only at close-out. (§6)
 26. Pin the import surface: after one real-tree run, require every module resolved under the repository to be covered by a hash pin, checked at the analyzer’s entry and again at its exit. (§6)
 27. Any pre-tag execution of the analyzer on the real tree is a disclosure event; log what it printed as knowledge held before the tag, in the design document, before the tag is cut. (§6)

The full record behind this paper, the design documents, the frozen analysis code, the 480 and 770 probe fits of the two campaigns, the campaign ledgers, and the adjudication rulings behind every number quoted here, sits under six git tags (`exp2-preregistered`, `exp2-closed`, `exp2b-preregistered`, `exp2b-closed`, `exp2c-preregistered`, `exp2c-closed`) in the supporting repository, github.com/lxman/decodable-is-not-learned: a history-preserving extraction of the working repository, public since the paper’s arXiv submission and archived at Zenodo under concept DOI 10.5281/zenodo.21830421, which resolves to the latest archived release. The same repository carries the successor experiments’ records under the same tag convention; Table C1 in Appendix C lists every tag, which of Section 6’s lessons each supplies, and the versioned DOI of every archived release. Interim stage tags exist only in the private repository, with the one exception of `exp1c-stage-a`, a power declaration made in advance. Commit SHAs quoted inside the ledgers resolve through that repository’s `PROVENANCE.md`.

Disclosure of AI assistance

This paper’s text was drafted by an AI assistant (Claude, Anthropic) working under my direction, and the same assistant operated the campaign infrastructure — run scripts, fit distribution, ledger transcription — across the experiments it reports. The division of labor is the one the campaign ledgers record by name: every design decision, gate ruling, capability ejection, and freeze decision is mine; the assistant executed and transcribed. Every number in the paper is transcribed from the tagged repository record or recomputed by committed scripts; every citation was verified against arXiv or the publisher’s record before entering the text; and I have read and approved every section. Responsibility for the content is mine.

(Every entry verified 2026-07-27: arXiv abstract pages fetched for each arXiv ID; venue records and the four pre-arXiv works checked against publisher pages. In-text years are arXiv submission years, i.e., first public availability; the published version is appended to each entry that has one.)

References

- Guillaume Alain and Yoshua Bengio. Understanding intermediate layers using linear classifier probes. *arXiv preprint arXiv:1610.01644*, 2016. Workshop version: ICLR 2017 Workshop Track.
- Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *arXiv preprint arXiv:2102.12452*, 2021. Journal version: Computational Linguistics, 48(1):207–219, 2022.
- Stella Biderman, Hailey Schoelkopf, Quentin Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, Aviya Skowron, Lintang Sutawika, and Oskar van der Wal. Pythia: A suite for analyzing large language models across training and scaling. *arXiv preprint arXiv:2304.01373*, 2023. Conference version: ICML 2023, PMLR 202:2397–2430.
- Thomas M. Cover. Geometrical and statistical properties of systems of linear inequalities with applications in pattern recognition. *IEEE Transactions on Electronic Computers*, EC-14:326–334, 1965.
- Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav Goldberg. Amnesic probing: Behavioral explanation with amnesic counterfactuals. *arXiv preprint arXiv:2006.00995*, 2020. Journal version: Transactions of the Association for Computational Linguistics, 9:160–175, 2021.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. *arXiv preprint arXiv:1909.03368*, 2019. Conference version: EMNLP-IJCNLP 2019, 2733–2743.
- Herbert Jaeger. The “echo state” approach to analysing and training recurrent neural networks. Technical Report GMD Report 148, German National Research Center for Information Technology, 2001.
- Wolfgang Maass, Thomas Natschläger, and Henry Markram. Real-time computing without stable states: A new framework for neural computation based on perturbations. *Neural Computation*, 14(11):2531–2560, 2002.
- Ali Rahimi and Benjamin Recht. Random features for large-scale kernel machines. In *Advances in Neural Information Processing Systems 20*, 2007.
- Rylan Schaeffer, Brando Miranda, and Sanmi Koyejo. Are emergent abilities of large language models a mirage? *arXiv preprint arXiv:2304.15004*, 2023. Conference version: NeurIPS 2023.
- Elena Voita and Ivan Titov. Information-theoretic probing with minimum description length. *arXiv preprint arXiv:2003.12298*, 2020. Conference version: EMNLP 2020, 183–196.
- Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, Ed H. Chi, Tatsunori Hashimoto, Oriol Vinyals, Percy Liang, Jeff Dean, and William Fedus. Emergent abilities of large language models. *arXiv preprint arXiv:2206.07682*, 2022. Journal version: Transactions on Machine Learning Research, 2022.

John Wieting and Douwe Kiela. No training required: Exploring random encoders for sentence classification. *arXiv preprint arXiv:1901.10444*, 2019. Conference version: ICLR 2019.

Kelly W. Zhang and Samuel R. Bowman. Language modeling teaches you more syntax than translation does: Lessons learned through auxiliary task analysis. *arXiv preprint arXiv:1809.10040*, 2018. Workshop version: BlackboxNLP at EMNLP 2018, 359–361.

A Per-capability detail

One example item per capability (question, free-form answer, and the probe label the instrument actually scores), then the complete starved-margin record: five seeds per cell, trained (m3) and untrained (known_absent), both sizes, transcribed by committed script from the fit files under tag `exp2b-closed`. Attrited capabilities first, alphabetical within fate.

A.1 One item per capability

- **bin2dec** — “What is binary 1111000111 in decimal?” → 967 (probe label: 7)
- **collatz2** — “Apply the Collatz rule (if even, halve; if odd, triple and add 1) twice to 4543. What is the result?” → 6815 (probe label: 0)
- **digitprod7** — “What is the product of the digits of 5545, mod 7?” → 3 (probe label: 3)
- **isqrt** — “What is the integer square root (rounded down) of 9470?” → 97 (probe label: 7)
- **mod7_add** — “What is $(57 + 38) \bmod 7$?” → 4 (probe label: 1)
- **mod7_mul** — “What is $(18 * 76) \bmod 7$?” → 3 (probe label: 4)
- **mul3x1** — “What is $267 * 4$?” → 1068 (probe label: 6)
- **numletter** — “Multiply 168 by the alphabet position of ‘j’. What is the result mod 26?” → 16 (probe label: 16)
- **roman** — “What is LXVII + LXXXVII in decimal?” → 154 (probe label: 7)
- **sq_mod7** — “What is $(22^2 + 99) \bmod 7$?” → 2 (probe label: 1)
- **units** — “How many milliliters are in 254 centiliters?” → 2540 (probe label: 1)
- **unscramble** — “Unscramble the letters ‘gloan’ to form an English word.” → along (probe label: a)
- **weekday** — “What day of the week is 252 days after Sunday?” → Sunday (probe label: 0)
- **add3_mid** — “What is $157 + 456$?” → 613 (probe label: 1)
- **add_base8** — “What is $67 + 76$ in base 8 (both numbers are octal)?” → 165 (probe label: 5)
- **antonym** — “Which of these means the opposite of ‘innocent’: busy, firm, exact, guilty?” → guilty (probe label: 4)
- **base7** — “Write 8622 in base 7.” → 34065 (probe label: 5)
- **caesar** — “The word ‘jsywd’ was made by shifting each letter of a word forward by 5. What was the original word?” → entry (probe label: e)
- **clock24** — “It is 2:00 on a 24-hour clock. What time will it be in 130 hours? Answer as H:00.” → 12:00 (probe label: 12)
- **count_div7** — “How many integers between 88 and 134 (inclusive) are divisible by 7?” → 7 (probe label: 7)
- **mod13** — “What is $(659 + 800) \bmod 13$?” → 3 (probe label: 9)
- **oct2dec** — “What is octal 6302 in decimal?” → 3266 (probe label: 6)
- **odd_one_out** — “Which word is not like the others: heron, swan, pear, owl?” → pear (probe label: 3)
- **reverse_string** — “Spell the string ‘vuldr’ backwards.” → rdluv (probe label: r)
- **sub3_mid** — “What is $328 - 124$?” → 204 (probe label: 0)

A.2 All starved margins, per seed

capability	410M untrained	410M trained	1B untrained	1B trained
bin2dec	0.08 / 0.08 / 0.06 / 0.00 / 0.06	0.00 / 0.07 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.07	0.00 / 0.00 / 0.00 / 0.00 / 0.00
collatz2	0.76 / 0.77 / 0.78 / 0.75 / 0.74	0.94 / 0.93 / 0.93 / 0.92 / 0.93	0.81 / 0.81 / 0.79 / 0.78 / 0.75	0.96 / 0.96 / 0.95 / 0.97 / 0.97
digitprod7	0.22 / 0.21 / 0.28 / 0.20 / 0.24	0.28 / 0.27 / 0.36 / 0.26 / 0.29	0.22 / 0.21 / 0.30 / 0.19 / 0.26	0.28 / 0.27 / 0.34 / 0.26 / 0.29
isqrt	0.41 / 0.36 / 0.38 / 0.42 / 0.38	0.62 / 0.57 / 0.62 / 0.63 / 0.57	0.44 / 0.37 / 0.43 / 0.40 / 0.38	0.66 / 0.63 / 0.67 / 0.66 / 0.63
mod7_add	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.10 / 0.00 / 0.09 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00
mod7_mul	0.00 / 0.10 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.11 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00
mul3x1	0.09 / 0.08 / 0.00 / 0.11 / 0.00	0.13 / 0.13 / 0.13 / 0.14 / 0.14	0.09 / 0.08 / 0.07 / 0.11 / 0.09	0.21 / 0.17 / 0.19 / 0.16 / 0.17
numletter	0.06 / 0.09 / 0.08 / 0.06 / 0.10	0.11 / 0.10 / 0.11 / 0.10 / 0.10	0.07 / 0.07 / 0.08 / 0.06 / 0.07	0.09 / 0.12 / 0.12 / 0.11 / 0.09
roman	0.65 / 0.77 / 0.68 / 0.66 / 0.78	0.84 / 0.86 / 0.82 / 0.89 / 0.96	0.64 / 0.74 / 0.69 / 0.72 / 0.82	0.89 / 0.90 / 0.86 / 0.84 / 0.90
sq_mod7	0.00 / 0.13 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.26 / 0.00 / 0.00 / 0.00	0.00 / 0.13 / 0.00 / 0.00 / 0.00
units	0.00 / 0.00 / 0.39 / 0.00 / 0.00	0.00 / 0.45 / 0.52 / 0.00 / 0.00	0.00 / 0.00 / 0.52 / 0.00 / 0.00	0.00 / 0.00 / 0.52 / 0.15 / 0.00
unscramble	0.20 / 0.18 / 0.22 / 0.18 / 0.20	0.32 / 0.25 / 0.29 / 0.31 / 0.33	0.21 / 0.19 / 0.22 / 0.21 / 0.19	0.33 / 0.30 / 0.32 / 0.32 / 0.33
weekday	0.00 / 0.09 / 0.00 / 0.11 / 0.00	0.09 / 0.08 / 0.00 / 0.00 / 0.00	0.08 / 0.12 / 0.09 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.09
add3_mid	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00
add_base8	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00
antonym	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.57 / 0.58 / 0.62 / 0.53 / 0.59	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.52 / 0.47 / 0.44 / 0.36 / 0.45
base7	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00
caesar	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00
clock24	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.05 / 0.00
count_div7	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.11 / 0.12 / 0.13 / 0.10 / 0.11	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.10 / 0.12 / 0.13 / 0.10 / 0.10
mod13	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00
oct2dec	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00
odd_one_out	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.33 / 0.25 / 0.33 / 0.32 / 0.39	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.24 / 0.23 / 0.28 / 0.31 / 0.30
reverse_string	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.57 / 0.52 / 0.62 / 0.60 / 0.55	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.69 / 0.58 / 0.70 / 0.78 / 0.61
sub3_mid	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00	0.00 / 0.00 / 0.00 / 0.00 / 0.00

B The frozen report and the gate scoreboard

The frozen report ran automatically at campaign end on 2026-07-27 and exited “GATE ATTENTION REQUIRED.” Its formal outputs, with the closeout adjudication of each (rulings a–d, all accepted, ledgered):

Gate 1, untrained weights. 0 tolerated floor fires; 86 structural leak fires across 250 fits, spanning 13 capabilities; pooled count test $p = 6.5e-117$ as coded. Ruled: attrition $\times 13$ per the design (the pooled abort is the implementation deviation of Section 6). Attrition list: bin2dec, collatz2, digitprod7, isqrt, mod7_add, mod7_mul, mul3x1, numletter, roman, sq_mod7, units, unscramble, weekday.

Gate 2, known-present. Entity tracking detected by seed majority at both sizes, mean starved margins .330 (410M) and .281 (1B); the copy control at .997/.998. PASS.

Gate 3, shuffled labels. 2 fires in 250 fits (410M only, zero at 1B), count test $p = .538$ against the floor rate; both fires failed the per-fire floor-signature predicate (3.6 and 4.7 null SD), which as coded projected a pipeline abort. Ruled: clean by count; the predicate is the calibration misspecification of Section 6; no abort.

Gate 4, argmax control. Copy control generated correctly at 0.868 at 410M (CP95 [.835, .896], upper bound below the frozen .9 bar): FAIL. 0.954 at 1B: PASS. Ruled a design-assumption finding (reliability measured on the prior battery at .994 did not transfer); documented caveat on the 410M inclusion readings.

Verdict. Twelve survivors against a floor of twenty: INSUFFICIENT_DATA. No Stage 1 was assembled; no eval-side model was queried. The verdict was projected in the ledger on 2026-07-23/24, before the report ran, and the report confirmed it.

One number in the closeout artifacts is superseded: they quote the ledger’s expectation of “~0.9 predicate-failing fires per 250” for a clean null, where the exact value is 1.6 (post-closeout correction, ledgered 2026-07-27; Section 6 carries the exact arithmetic).

C Reproduction notes

Models. EleutherAI Pythia, standard (non-deduped) branch, final checkpoints, fp16: pythia-410m at revision 9879c9b5f8bea9051dcb0e68dff21493d67e9d4f, pythia-1b at f73d7dcc545c8bd326d8559c8ef84ffe92fea6b2 (main-branch SHAs resolved and ledgered 2026-07-07; run code loads by SHA, so an upstream branch move cannot swap weights silently). The untrained twin at each size is the same configuration instantiated under `torch.manual_seed(0)` with no pretrained weights. The design named 2.8b, 6.9b, and 12b as eval-side models; none was ever queried.

Probe. Standard scaler + logistic regression ($C = 1.0$, `max_iter = 100`), 2,500 label permutations under the fixed split, Bonferroni families of 18 (410M) and 14 (1B) (layer, position) candidates, corrected $\alpha = .01$; probe seeds 0–4, each redrawing the starved holdout. All thresholds frozen at tag `exp2b-preregistered`.

Environment. Python 3.11 (Homebrew), torch 2.12.1, transformers 5.13.0; activation collection on Apple silicon (MPS, fp16); probe fits distributed across arm64 Accelerate, x86-64 OpenBLAS, and AMD64 Windows under single-threaded BLAS, each machine admitted only after reproducing the reference fixture bit for bit through the real probe path.

A later re-fit of the whole campaign from its stored activations found that all 731 fits computed on the reference machine reproduce their stored accuracy exactly, while 15 of the 39 computed on the two remote workers, both of which had passed the fixture gate, do not, by at most 0.0098 and a median of 0.0026, one to four items on these validation sets: a different BLAS accumulating a different sum, moving no verdict and no taxonomy assignment. One fixture, exercising one code path on one input, cannot certify that every subsequent fit on that machine lands on the same floating-point value. The rule it promotes is checklist item 13: certify a worker on a random sample of the units it actually produced, and where exact cross-machine agreement is load-bearing, compute those fits on one machine.

Artifacts. 770 probe-fit JSONs (Experiment 2b) and 480 (Experiment 2) under the experiments’ results trees; `m2_report.json`; campaign and inclusion logs; SHA-256 digests of all 108 activation files (`activations_sha256.txt`; the npz files stay local by convention). Battery item files carry per-capability

generation seeds. Figures 2 and 3, Table 1, and Appendix A regenerate from these artifacts via the committed scripts in `paper/`. The tags `exp2-preregistered`, `exp2-closed`, `exp2b-preregistered` and `exp2b-closed` mark the freeze and closeout commits of the two campaigns, and `exp2c-preregistered` the successor battery’s freeze (Section 7). All of it is in the supporting repository (Section 10).

The record’s tags and releases. Table C1 lists the supporting repository’s experiment tags and the versioned Zenodo releases; the concept DOI in Section 10 resolves to the latest.

Table C1. Experiment tags in the supporting repository. Each experiment carries a `-preregistered` tag at its freeze and a `-closed` tag at its verdict; the sealed-stage tags are named where they exist.

experiment	tags	in this paper
2, 2b, 2c	<code>exp2-*</code> , <code>exp2b-*</code> , <code>exp2c-*</code>	Sections 3–8
3a, 3b, 3	<code>exp3a-*</code> , <code>exp3b-*</code> , <code>exp3-*</code>	§8 (3b); §6 freeze block (3)
3c, 3d, 3e	<code>exp3c-*</code> , <code>exp3d-*</code> , <code>exp3e-*</code>	§6 lessons 5, 6, 7
2d, 2e	<code>exp2d-*</code> , <code>exp2e-*</code>	§6 lesson 8 (2d); 2e named, not drawn on
2f	<code>exp2f-*</code>	§6 lesson 9
2g, 2h	<code>exp2g-*</code> (+ <code>-predictor-sealed</code>), <code>exp2h-*</code>	named, not drawn on
2i	<code>exp2i-*</code> (+ <code>-predictor-sealed</code> , <code>-endpoint-sealed</code> ; the preregistration tag re-cut after a campaign stop, disclosed in PROVENANCE)	§6 lesson 10
2j	<code>exp2j-*</code>	§6 lesson 11
2k	<code>exp2k-*</code> (+ <code>-predictor-sealed</code>)	named, not drawn on
2l, 2m, 2n	<code>exp2l-*</code> , <code>exp2m-*</code> , <code>exp2n-*</code> (each + <code>-endpoint-sealed</code>)	named, not drawn on
4	<code>exp4-*</code> (+ <code>-reference-sealed</code> ; the preregistration tag re-cut after a campaign stop, disclosed in PROVENANCE)	named, not drawn on
4b	<code>exp4b-*</code>	named, not drawn on
4c	<code>exp4c-*</code> (the preregistration tag re-cut after a preflight amendment, disclosed in PROVENANCE)	named, not drawn on
5	<code>exp5-*</code> (+ <code>-targets-sealed</code>)	named, not drawn on
5b	<code>exp5b-*</code>	named, not drawn on
1, 1b, 1c	<code>exp1-analysis-frozen</code> , <code>exp1b-*</code> , <code>exp1c-*</code> (+ <code>exp1c-stage-a</code>)	companion essay; not drawn on

Archived releases (all 10.5281/zenodo. prefixed): v1.0 21830422 (the original snapshot); v1.1 21998671 (Experiments 3 and 3c); v1.2 22011547 (3d); v1.3 22045940 (3e); v1.4 22056230 (2d); v1.5 22059612 (2e); v1.6 22063906 (2f); v1.7 22064573 (1, 1b, 1c); v1.8 22064649 (this paper’s ninth lesson); v1.9 22087512 (2g); v1.10 22104580 (2h); v1.11 22145025 (2i); v1.12 22151832 (2j); v1.13 22151979 (this paper’s eleventh lesson); v1.14 22177233 (2k); v1.15 22286867 (2l); v1.16 22413760 (2m); v1.17 22667987 (2n); v1.18 22793373 (4); v1.19 22819695 (4b); v1.20 22901051 (4c); v1.21 22962602 (5); v1.22 22998653 (5b).